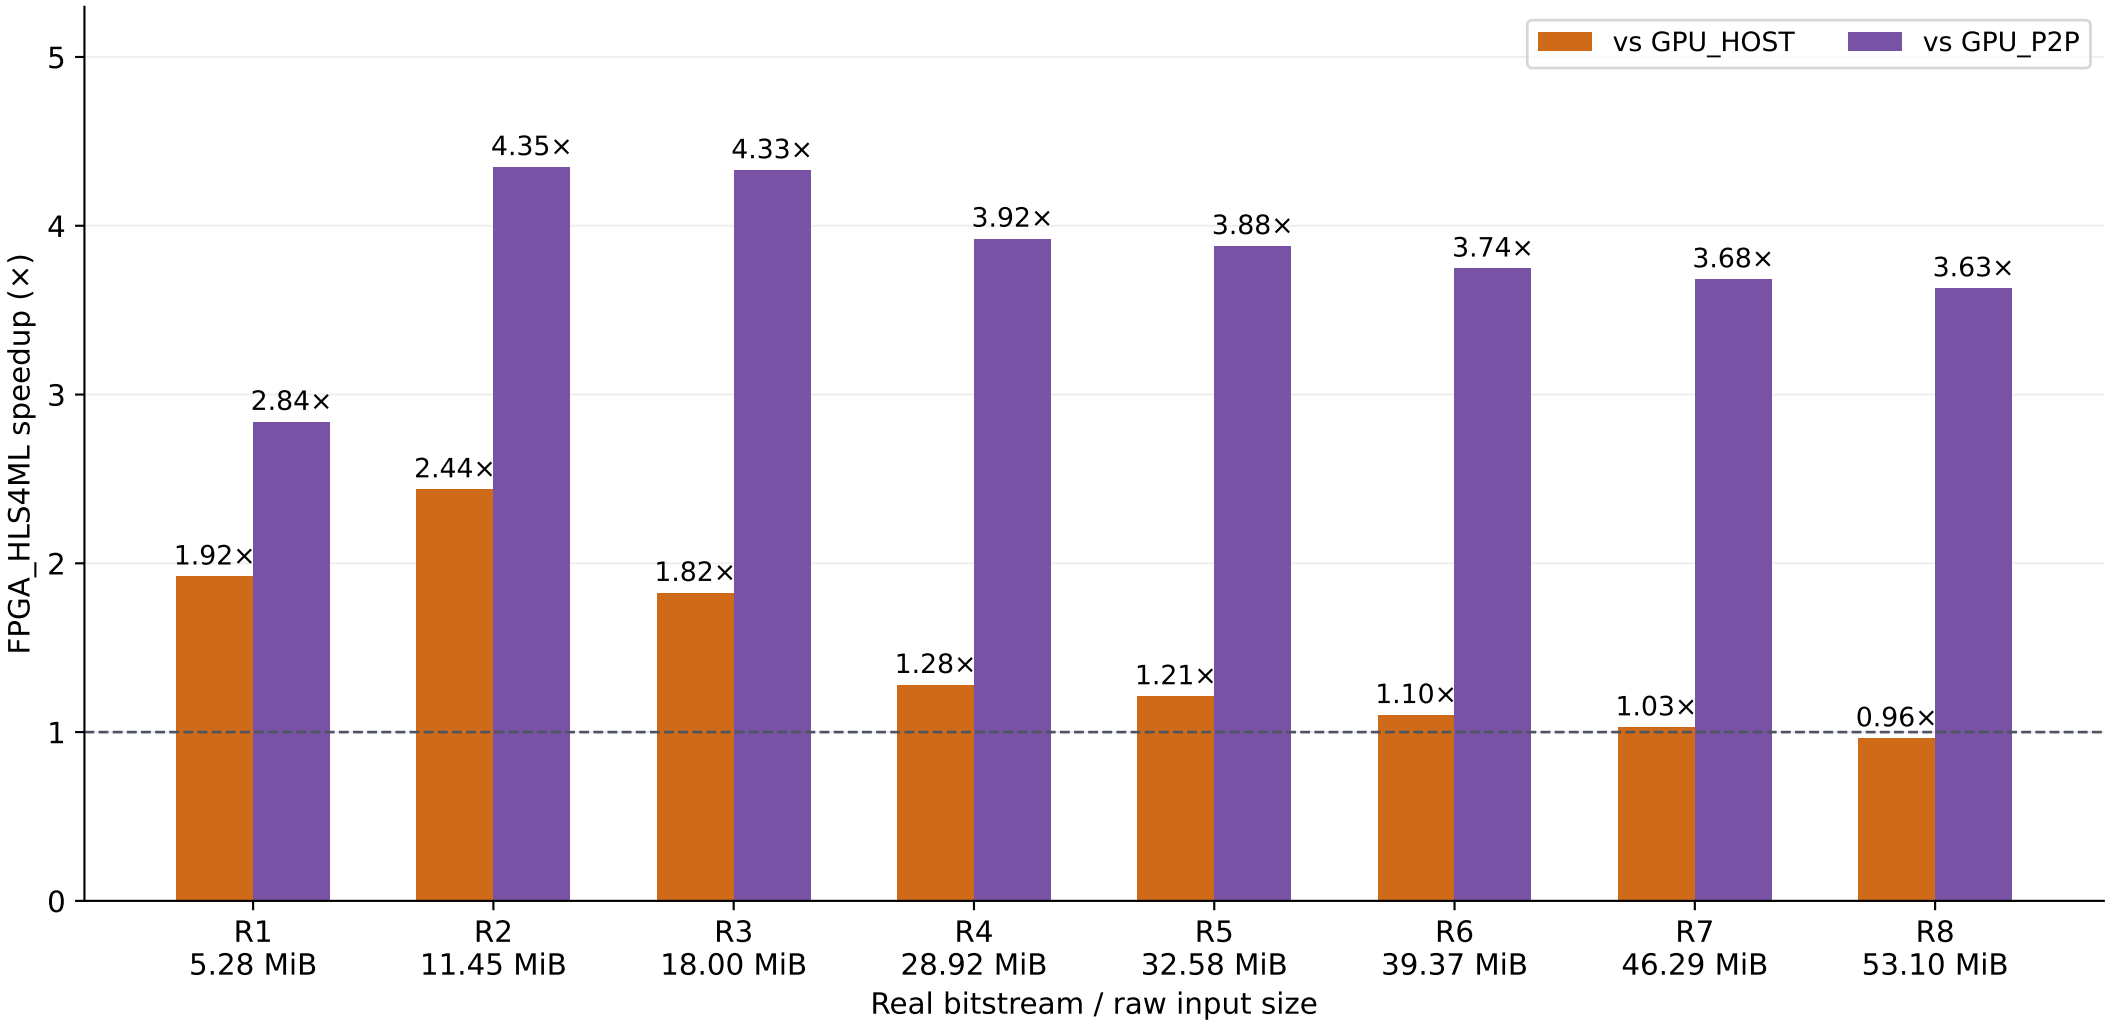
Real bitstreams: FPGA_HLS4ML speedup to verdict availability
Speedup = GPU median / FPGA_HLS4ML median; above 1× favors FPGA_HLS4ML. Endpoint: FPGA result acceptance.
vs GPU_HOST
vs GPU_P2P
FPGA_HLS4ML speedup (×)
0
1
2
3
4
5
1.92×
2.84×
2.44×
4.35×
1.82×
4.33×
1.28×
3.92×
1.21×
3.88×
1.10×
3.74×
1.03×
3.68×
0.96×
3.63×
R1
5.28 MiB
R2
11.45 MiB
R3
18.00 MiB
R4
28.92 MiB
R5
32.58 MiB
R6
39.37 MiB
R7
46.29 MiB
R8
53.10 MiB
Real bitstream / raw input size