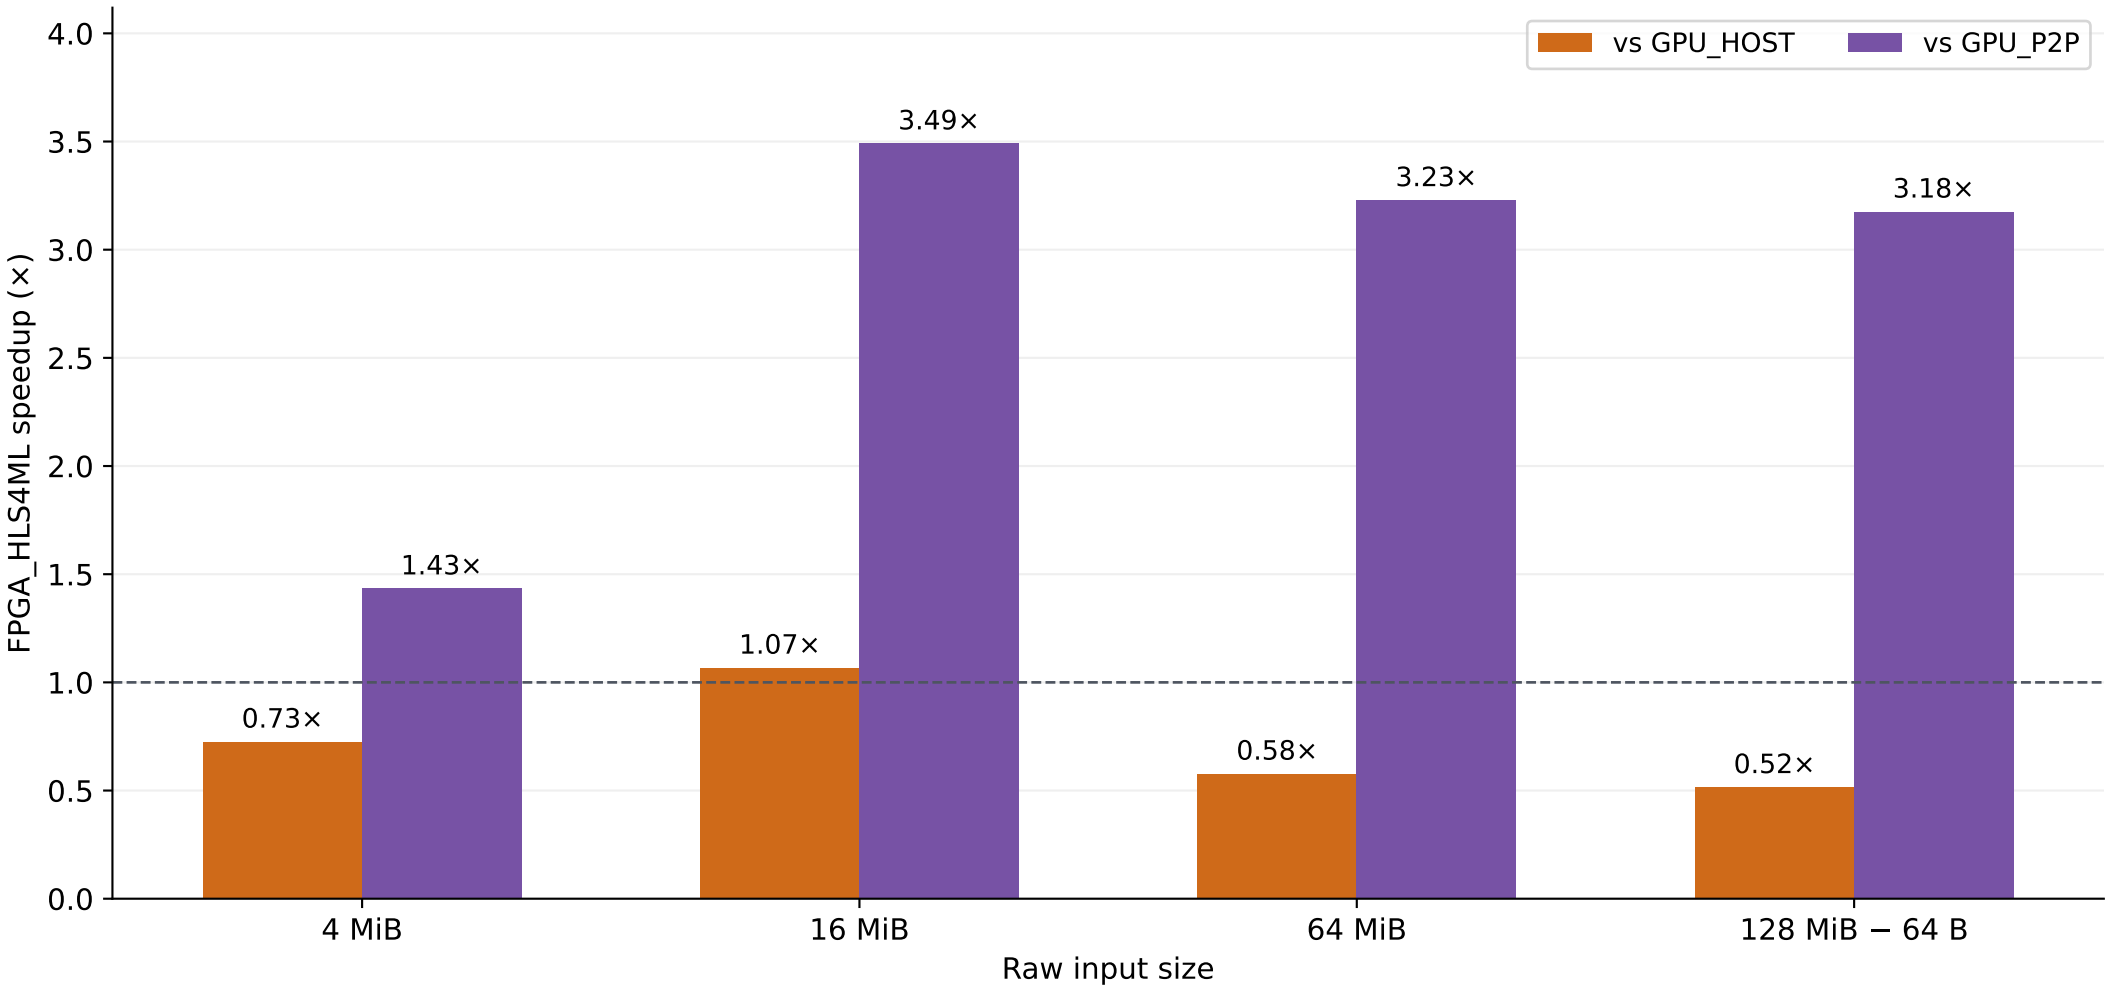
Synthetic inputs: FPGA_HLS4ML speedup to verdict availability
Speedup = GPU median / FPGA_HLS4ML median; above 1× favors FPGA_HLS4ML. Endpoint: FPGA result acceptance. GPU: R06 XLA + ConcreteFunction (FP32); FPGA: W8A8.
vs GPU_HOST
vs GPU_P2P
FPGA_HLS4ML speedup (×)
4.0
3.5
3.0
2.5
2.0
1.5
1.0
0.5
0.0
0.73×
1.43×
1.07×
3.49×
0.58×
3.23×
0.52×
3.18×
4 MiB
16 MiB
64 MiB
128 MiB − 64 B
Raw input size