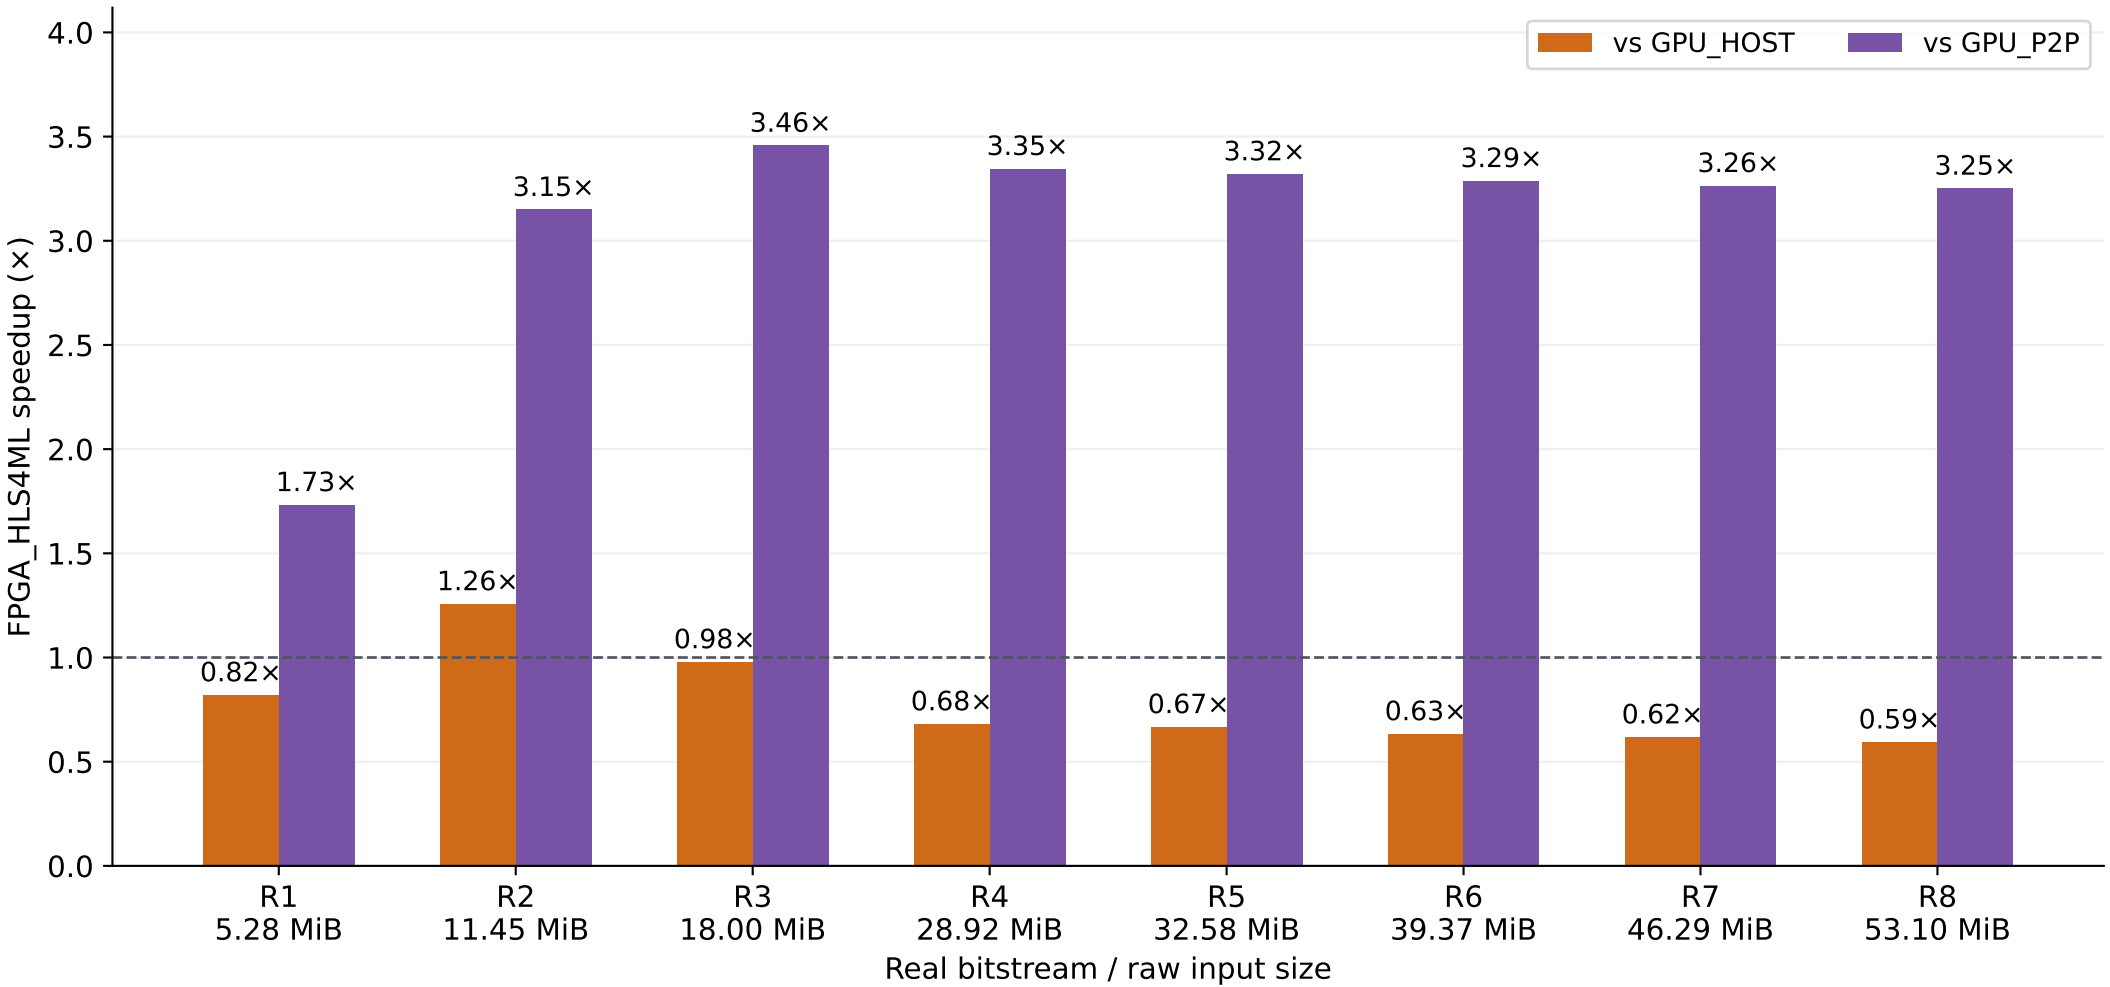
Real bitstreams: FPGA_HLS4ML speedup to verdict availability
Speedup = GPU median / FPGA_HLS4ML median; above 1× favors FPGA_HLS4ML. Endpoint: FPGA result acceptance. GPU: R06 XLA + ConcreteFunction (FP32); FPGA: W8A8.
vs GPU_HOST
vs GPU_P2P
FPGA_HLS4ML speedup (×)
4.0
3.5
3.0
2.5
2.0
1.5
1.0
0.5
0.0
0.82×
1.73×
1.26×
3.15×
0.98×
3.46×
0.68×
3.35×
0.67×
3.32×
0.63×
3.29×
0.62×
3.26×
0.59×
3.25×
R1
5.28 MiB
R2
11.45 MiB
R3
18.00 MiB
R4
28.92 MiB
R5
32.58 MiB
R6
39.37 MiB
R7
46.29 MiB
R8
53.10 MiB
Real bitstream / raw input size