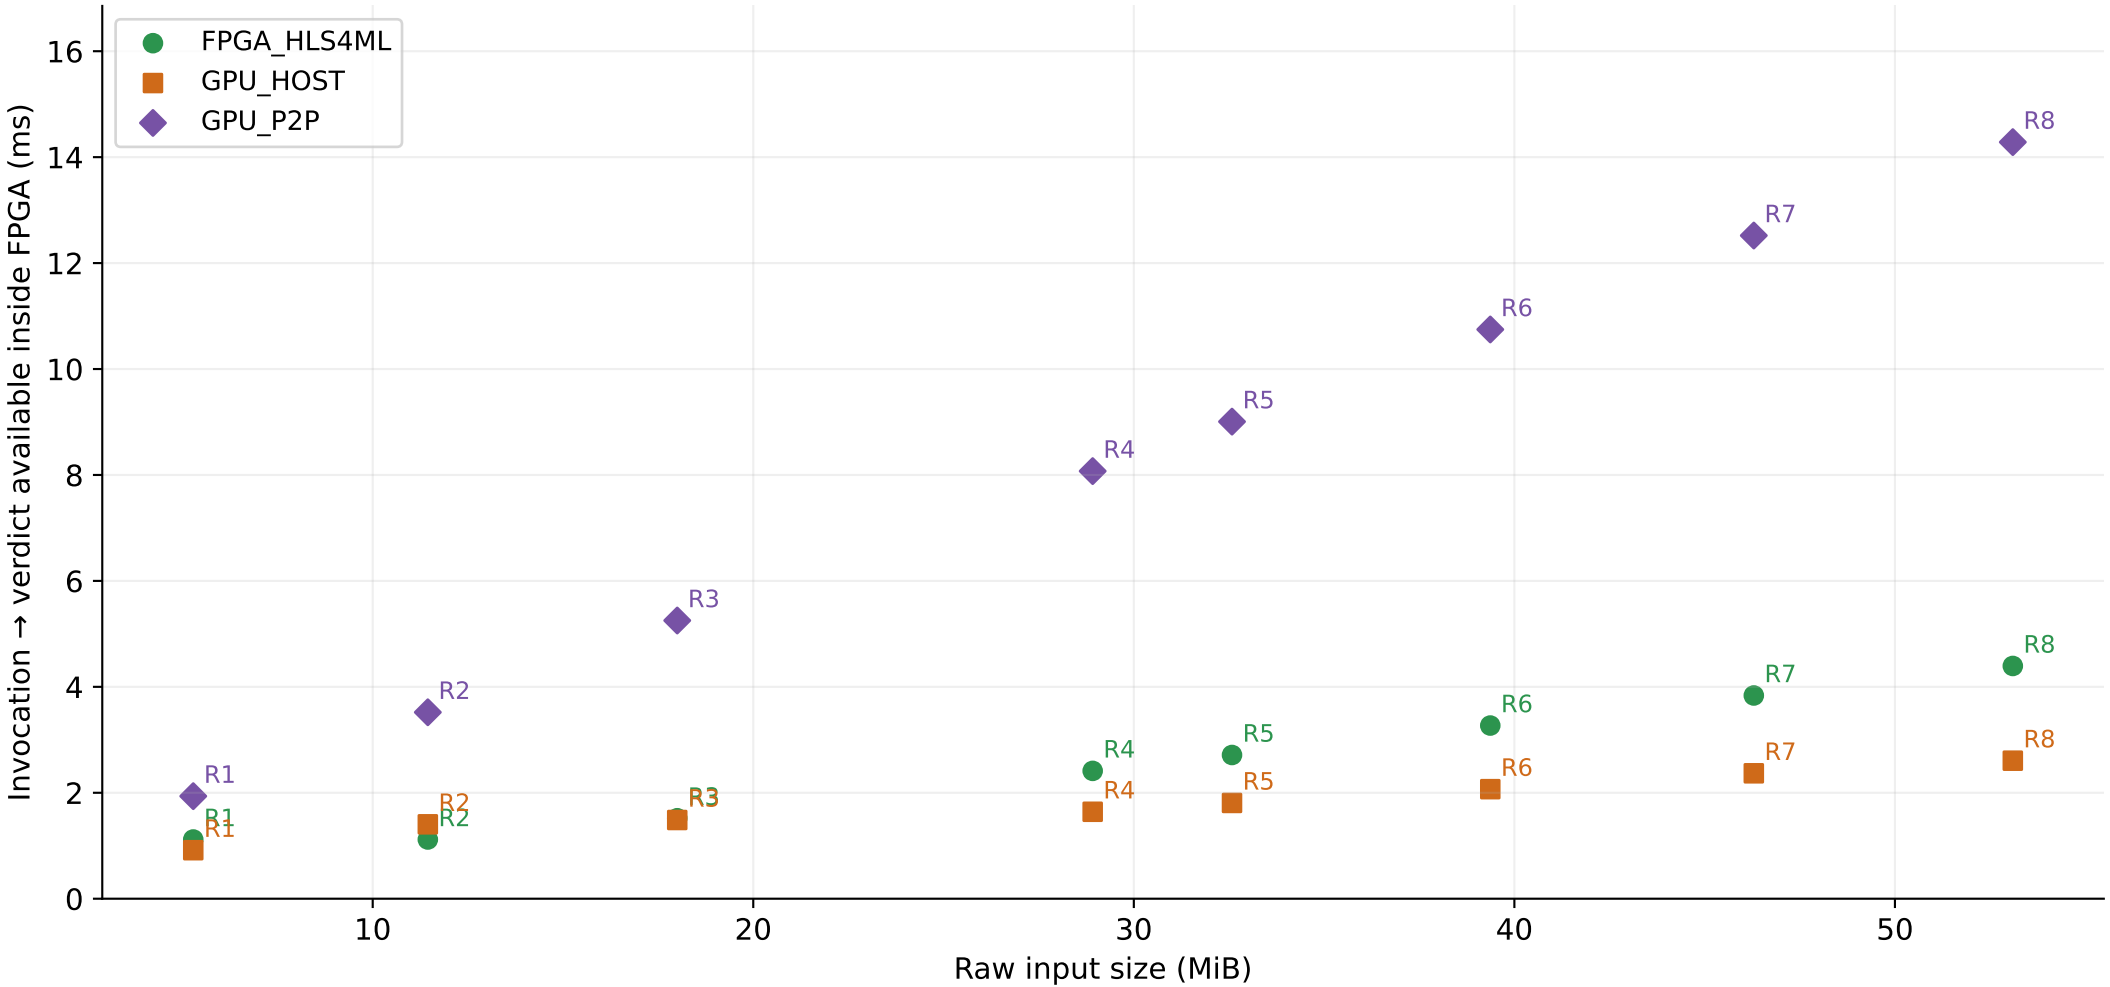
Real bitstreams: verdict-availability latency versus input size
Prepared host input → FPGA result acceptance. Points = medians of 200 calls per file/variant. GPU: R06 XLA + ConcreteFunction (FP32); FPGA: W8A8.
FPGA_HLS4ML
GPU_HOST
GPU_P2P
Invocation → verdict available inside FPGA (ms)
Raw input size (MiB)
0
2
4
6
8
10
12
14
16
10
20
30
40
50
R1
R2
R3
R4
R5
R6
R7
R8