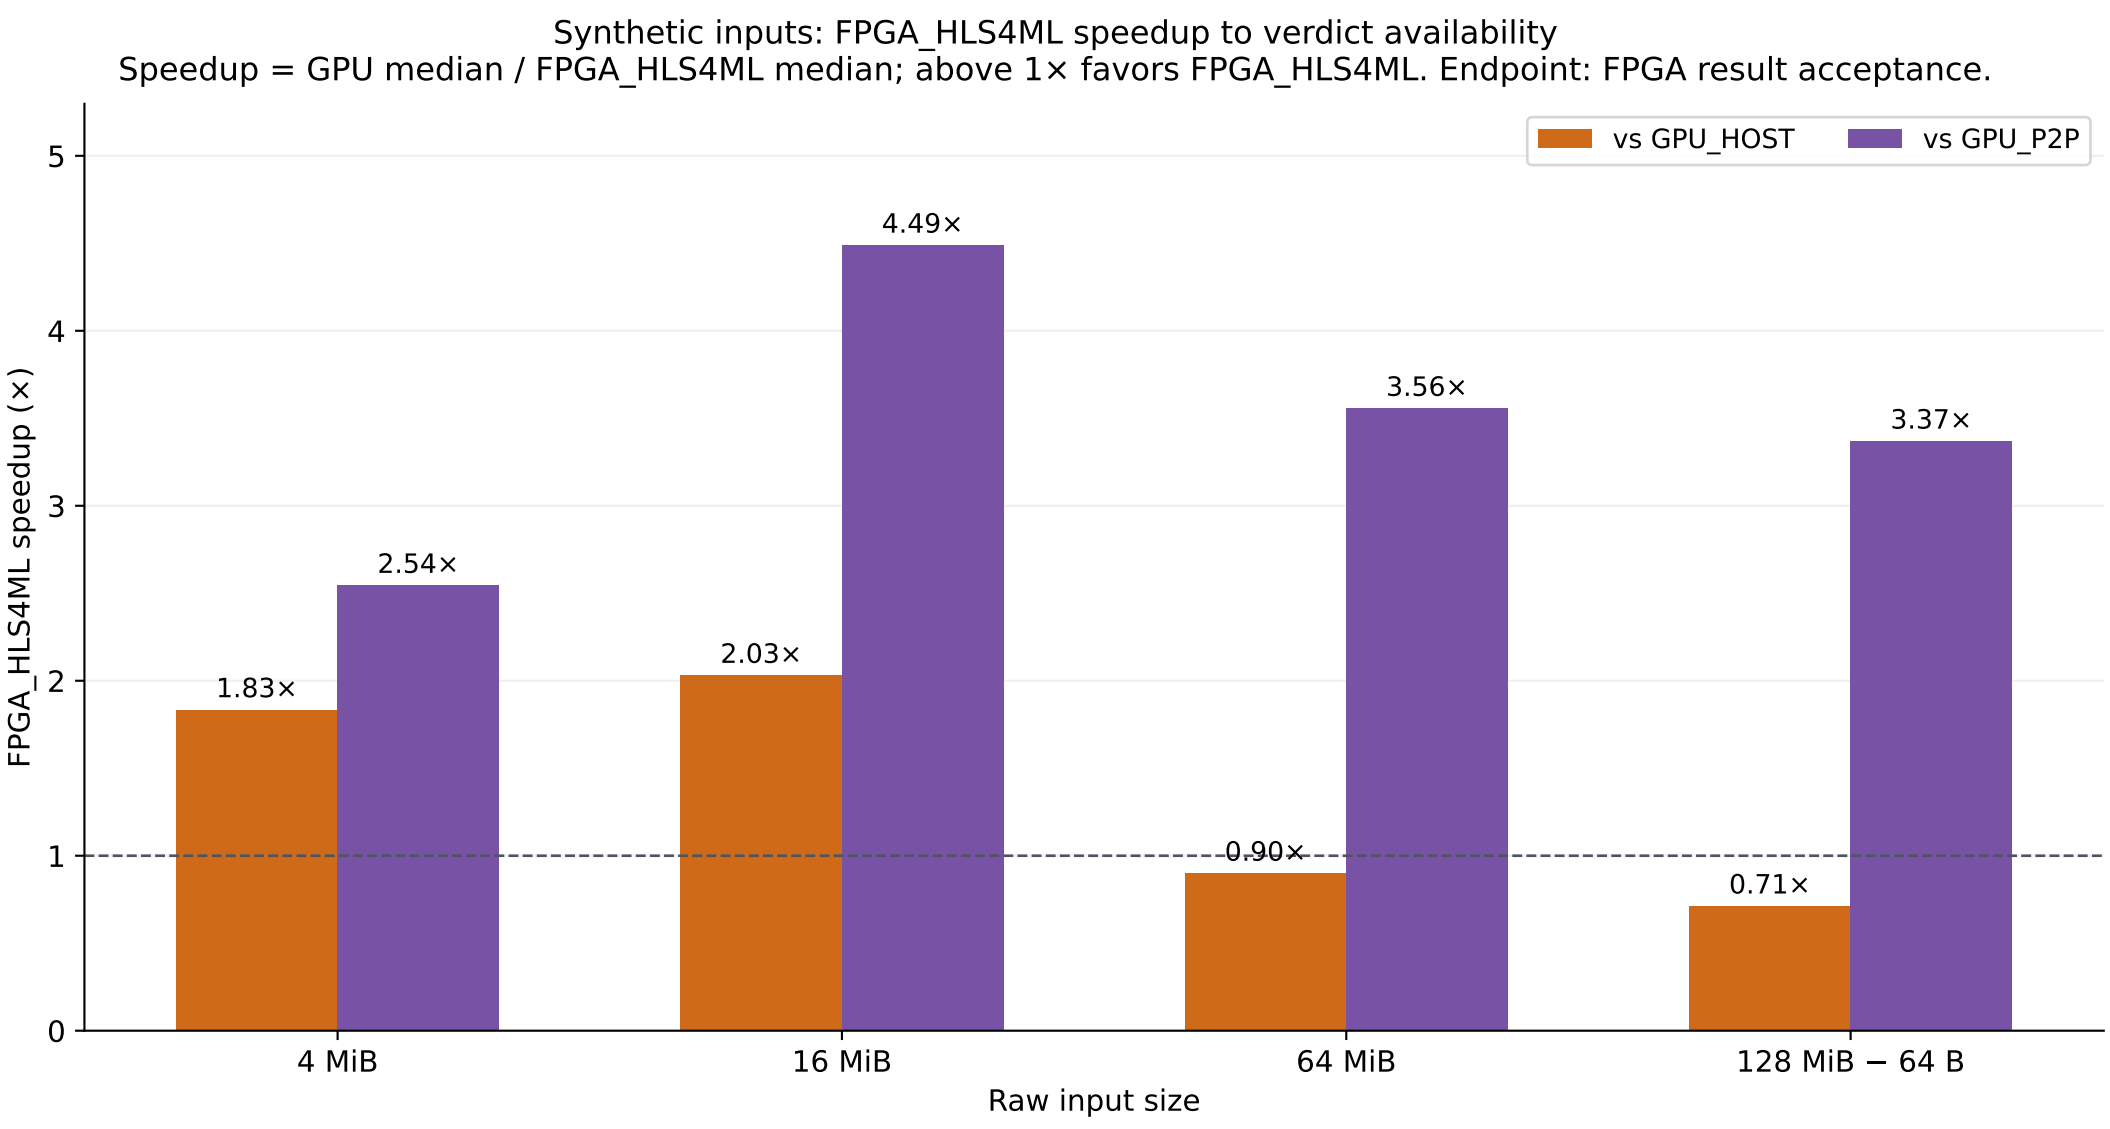
Synthetic inputs: FPGA_HLS4ML speedup to verdict availability
Speedup = GPU median / FPGA_HLS4ML median; above 1× favors FPGA_HLS4ML. Endpoint: FPGA result acceptance.
vs GPU_HOST
vs GPU_P2P
FPGA_HLS4ML speedup (×)
0
1
2
3
4
5
1.83×
2.54×
2.03×
4.49×
0.90×
3.56×
0.71×
3.37×
4 MiB
16 MiB
64 MiB
128 MiB − 64 B
Raw input size